



Introduction to HPC on Blanca

Introduction to HPC on Blanca

- Andrew Monaghan
- *Email:* Andrew.Monaghan@Colorado.edu
- *RC Homepage:* <https://www.colorado.edu/rc>

- Slides available for download at:
https://github.com/ResearchComputing/APPM_HPC

Outline for this presentation

- Part 1: Today:
 - Overview of CU Research Computing (CURC) and our resources
- Part 2: Next week
 - Using Blanca
 - Logging in
 - Basic Linux commands
 - File editing
 - Linux filesystem
 - Environment variables
 - Software modules on Blanca
 - Bash scripts and Job Scheduling

Part 1 (Overview)

What is Research Computing?

- Provide services for researchers that include:
 - High performance computing (HPC)
 - Data visualization
 - Data storage
 - High speed data transfer
 - Data management support
 - Consulting
 - Training
- We are likely best known for:
 - Alpine Supercomputer (~22,000 cores)
 - PetaLibrary storage
 - **Blanca "condo" cluster** (~4,000 cores)

High Performance Computing (HPC) vs. Traditional Computing

- Traditional computing *generally* has access to a single processor (perhaps multiple cores)

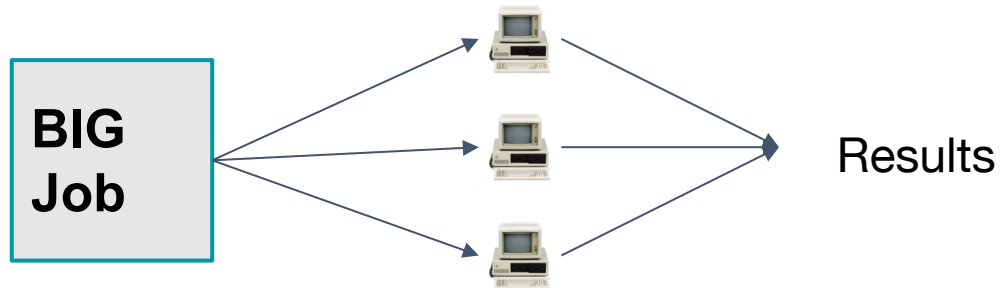


What can I use HPC for?

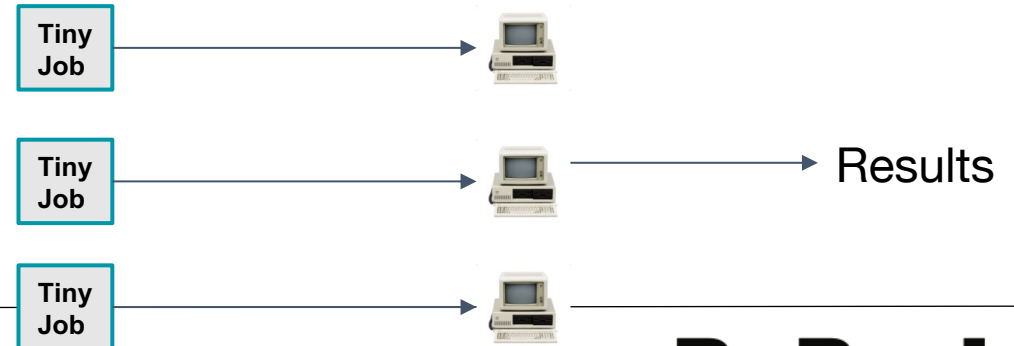
- Solving large problems that require more:
 - Memory than you have on your PC
 - cores/nodes/power than you have on your PC
- Jobs that require hardware you may not have:
 - High Performance GPU computing
 - Specific Operating System
- Visualization rendering

What can I use HPC for?

- Jobs that would take a long time on local machines can instead be distributed over hardware:
 - Parallelized to split up then joined (if software enabled)



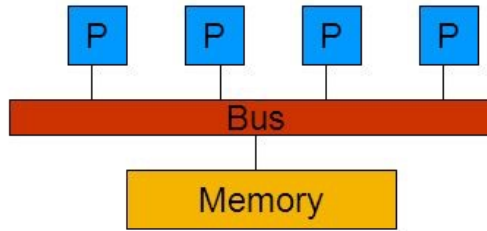
- Broken up into many serial jobs



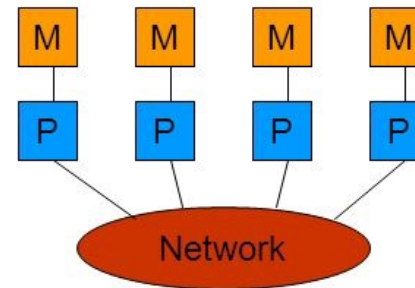
HPC Parallelization

HPC is set up for:

- shared memory (single node) parallelization
- distributed memory (multi-node) parallelization.



- Shared memory



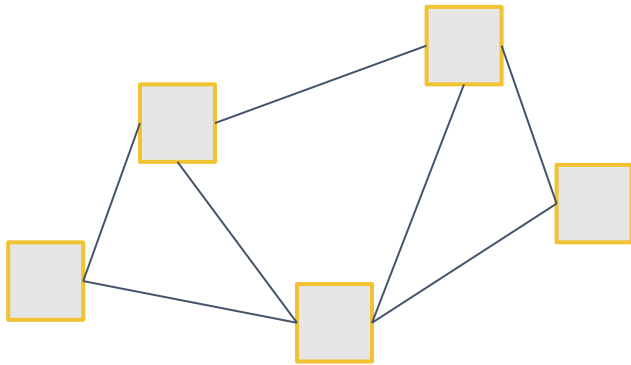
- Distributed memory

Source: https://images.slideplayer.com/25/7599921/slides/slide_4.jpg

Research Computing Resources

HPC Cluster: Alpine

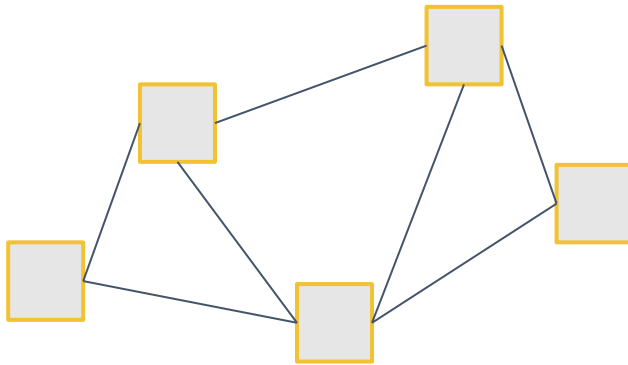
Alpine



- Alpine is the 3rd-generation HPC cluster at CURC, following:
 - Janus
 - RMACC Summit
- Alpine is a heterogeneous cluster with hardware currently provided by CU Boulder, CSU, and Anschutz
- Access available to CU Boulder, CSU, AMC and RMACC users

HPC Cluster: Alpine

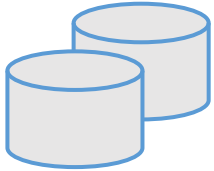
Alpine



- Hardware on Alpine will continue to be purchased and released in stages:
- Alpine (stage 3):
 - 184 General CPU Nodes
 - *AMD Milan, 64 Core, 3.74G RAM/Core*
 - 8 NVIDIA GPU Nodes
 - *3x NVIDIA A100 (atop General CPU node)*
 - 8 AMD GPU Nodes
 - *3x AMD MI100 (atop General CPU node)*
 - 12 AMD High-Memory Nodes
 - *AMD Milan, 48 Core, 21.5G RAM/Core*
 - Additional Hardware contributed by CSU, AMC
 - *Nodes which boost priority for CSU/AMC users*

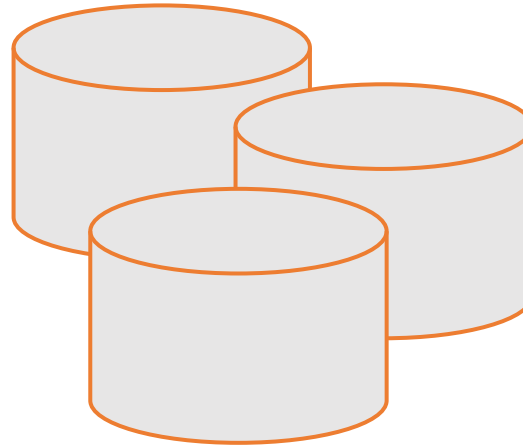
Storage at CURC

Core



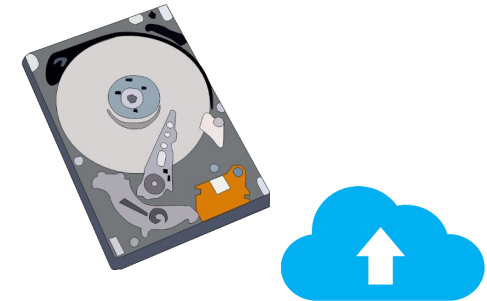
- Included with RC account
 - /home (2 GB/user)
 - /projects (250 GB/user)
 - scratch space (10 TB/user)

PetaLibrary



- Paid Service for:
 - Storage
 - Archive
 - Sharing of research data

Local or Cloud



- You can download your data locally or to a variety of other cloud resources
- Cloud Foundations at Research Computing

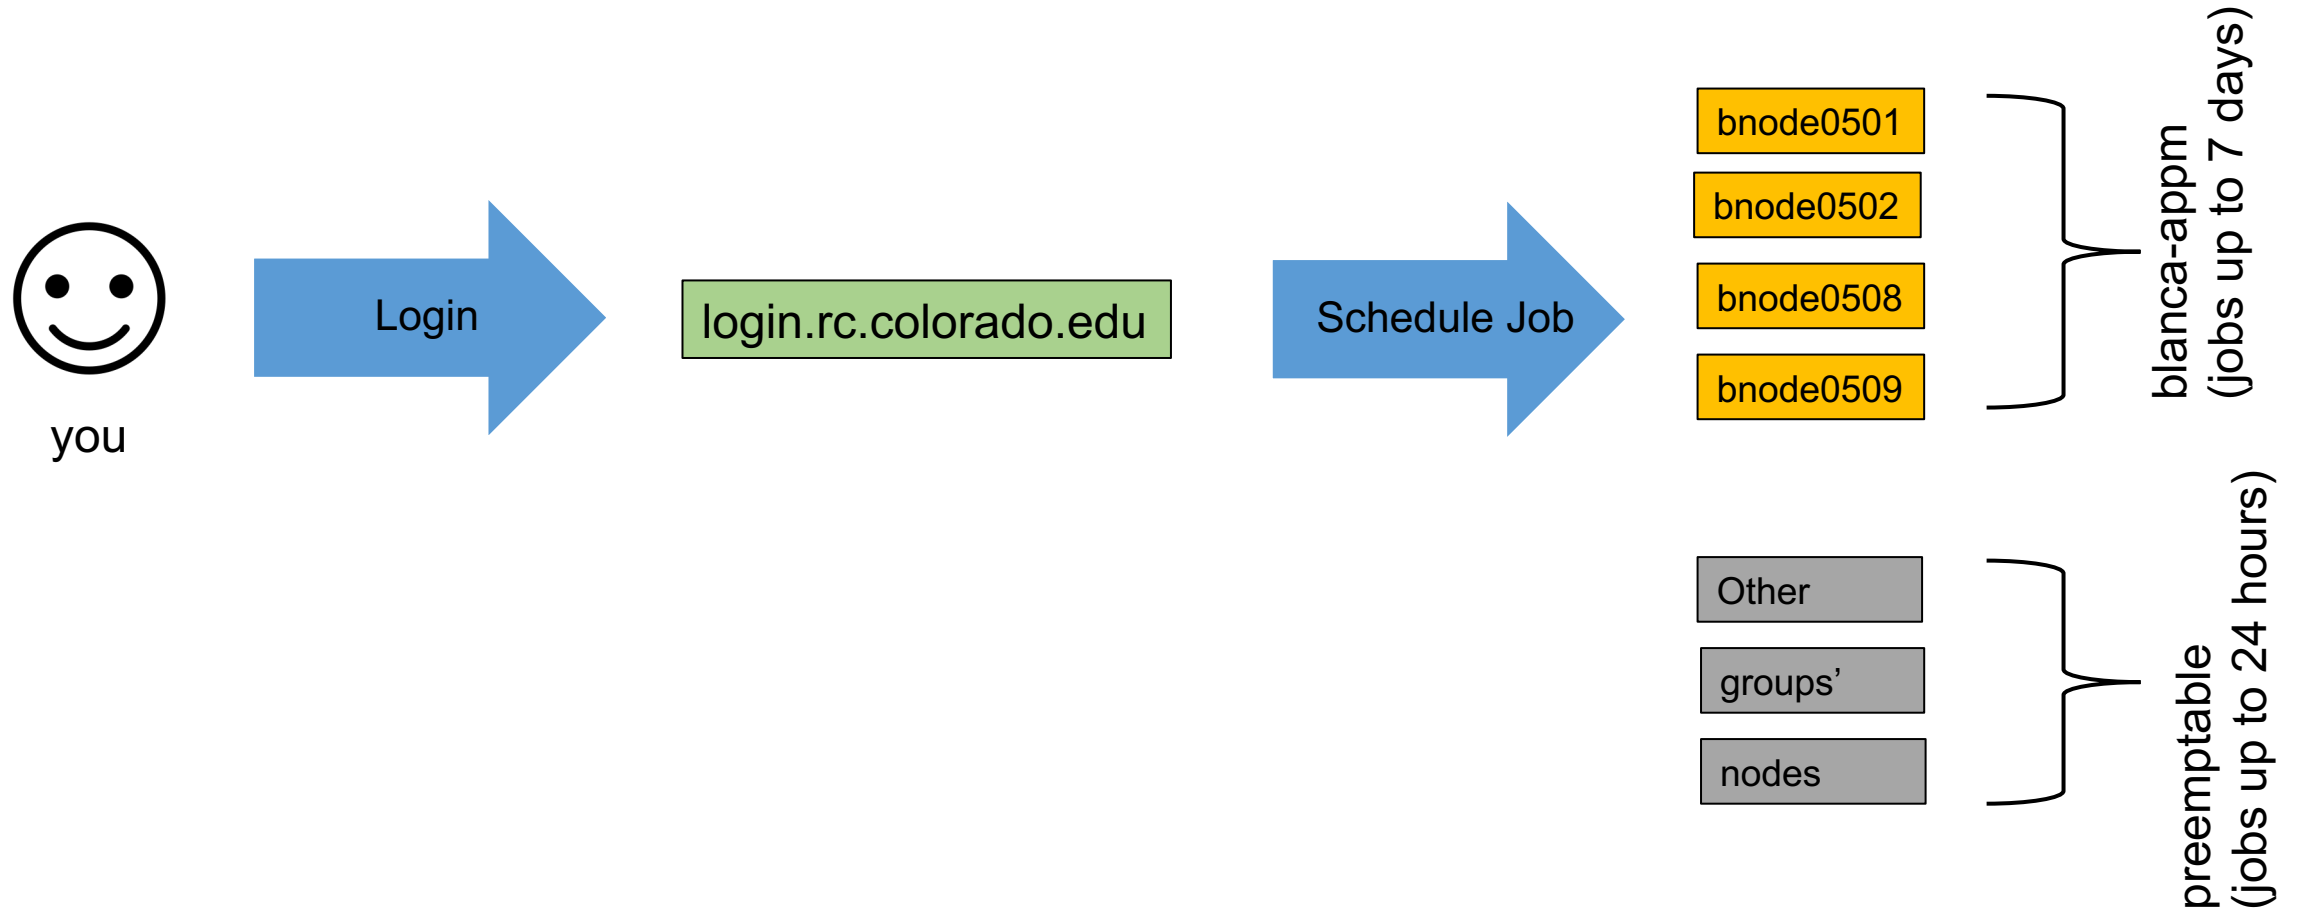
Blanca

- A “condo” cluster whereby individual research groups own nodes
- List of nodes and groups can be found [here](#)
- Users have dedicated access to their group’s nodes (e.g., blanca-appm)
 - Jobs up to 7 days long.
 - Can also run ‘preemptable’ jobs on other groups nodes (jobs up to 24 hours long)
- More documentation on Blanca:
<https://curc.readthedocs.io/en/latest/access/blanca.html>

Blanca APPM nodes

- bnode0501, bnode0502, (2 nodes)
 - 32 (effectively 64) cores, avx2, Cascade, 2.3 GB RAM/core)
- bnode0508, bnode0509 (2 nodes)
 - 40 (effectively 80) cores, avx2, Cascade, 2.3 GB RAM/core)

Blanca Workflow



Part 2 (Using Blanca)

Logging In

- `ssh <identkey>@login.rc.colorado.edu`
- Enter your `identkey_password`
- Authenticate by accepting the Duo push to your smartphone
 - Can also authenticate by text message, phone call, or token
- More info here:
<https://curc.readthedocs.io/en/latest/access/logging-in.html>

Linux

- Part of the Unix-like family of operating systems.
- Started in early '90s
- Several distributions are available – from enterprise-grade, like RedHat Linux (RHEL), to more consumer-focused, like Ubuntu.
 - Blanca nodes presently run RHEL7
- Runs on everything from embedded systems to supercomputers.
- Linux is simple, flexible, fast, many potent tools

Anatomy of a Linux command

- command [flags] [flag arguments] [target(s)]
 - `ls -l myworkdir/`
- Case is important!
- Help on commands is available through the "man" command (short for manual). E.g.,
 - `man ls`

File and directory related commands

- **pwd** – prints full path to current directory
- **cd** – changes directory; can use full or relative path as target
- **mkdir** – creates a subdirectory in the current directory
- **rm** – removes a file (`rm -r` removes a directory and all of its contents)
- **cp** – copies a file
- **mv** – moves (or renames) a file or directory
- **ls** – lists the contents of a directory (`ls -l` gives detailed listing)

File-viewing commands

- **more** – displays a file one screen at a time
- **cat** – prints entire file to the screen
- **head** – prints the first few lines of a file
- **tail** – prints the last few lines of a file

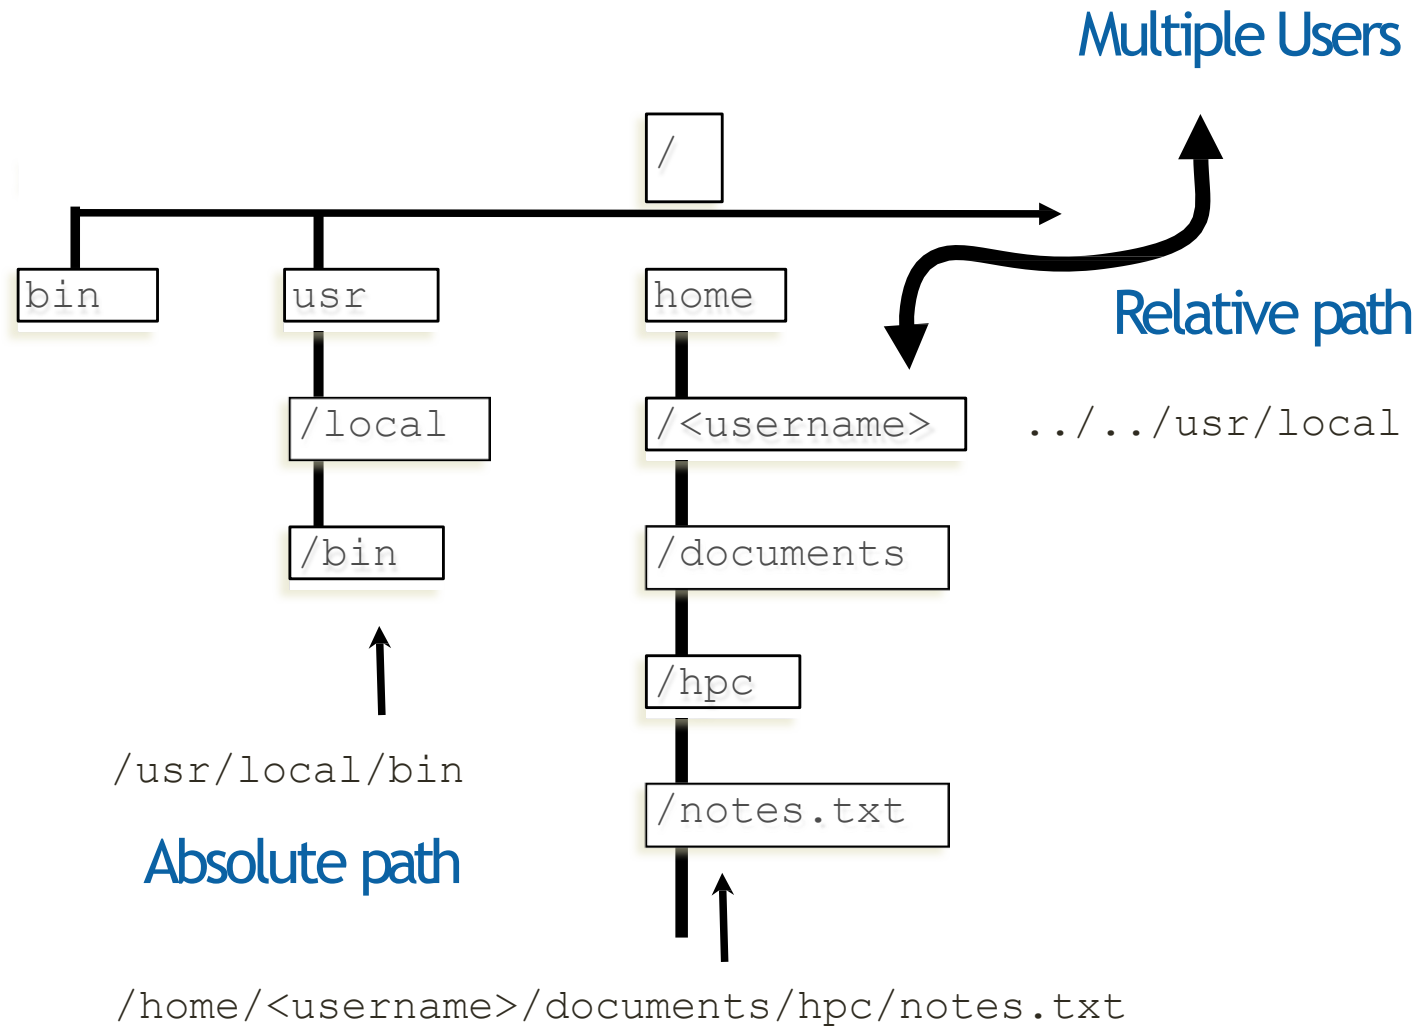
File editing with **nano**

- To edit a file:
 - `nano myfile.txt`
- From within Nano:
 - `Ctrl+o` save (need to confirm filename)
 - `Ctrl+x` exit
 - `Ctrl+k` cut
 - `Ctrl+u` paste
- Other popular Linux editors: vi, emacs

The Linux Filesystem

- System of arranging files on disk
- Consists of directories (folders) that can contain files or other directories
- Levels in full paths separated by *forward* slashes, e.g.
- `/home/user/scripts/analyze_data.sh`
- Case-sensitive; spaces in names discouraged
- Some shorthand:
 - . (the current directory)
 - .. (the directory one level above)
 - ~ (home directory)

Filesystem



Your personal directories on Blanca

- **/home/<username>**
 - Very small: 2GB.
 - Backed up daily.
 - Good for 'can't lose' files
- **/projects/<username>**
 - 250 GB
 - Backed up regularly
 - Good for storing scripts, self-installed software, some data
- **/rc_scratch/<username>**
 - Group-shared 130 TB partition.
 - Good for jobs with lots of I/O
 - Not backed up
 - Temporary: data deleted 90 days from creation.

Software

- Common software is available to everyone on the systems
- Research Computing uses modules to manage software
 - You load modules to prepare your environment for using software
 - Modules set any environment variables, paths, etc.
 - Set environment so application can find appropriate libraries, etc.
- You can also install your own software
 - It is best if you are responsible for support
 - We are happy to assist

Using Modules

- Must be on a **compute** node to browse the modules (e.g, bnode0501)
- To set up your environment to use a software package, type **module load <package>/<version>**
- Some modules might require a specific hierarchy to load
 - For some modules, you may need to specify a specific version
 - For example, **module load R/3.3.0**
 - For other modules, you may be able to be more generic
 - For example, **module load matlab**
- Some modules may require you to first load other modules that they depend on
- To find dependencies for a module, type **module spider <package>**
- To find out what software is available, you can type **module avail**

Using Anaconda for Python, R

- Many users prefer to use conda to manage python (and R, etc) packages and environments.
- We have Anaconda python installed on our system!
 - (but it isn't a module...)
- Documentation on use:
<https://curc.readthedocs.io/en/latest/software/python.html>

Next topic: Job scheduling

- **Job** – a system allotment of resources that run a particular application.
- **Slurm** -- resource manager
- Because Summit and Blanca are shared resources among a variety of groups on campus, users must run their applications through jobs.
 - Ensures everyone can utilize the system
 - No one person taking up too much of the system!

Two types of jobs

- **Interactive Jobs**

- Interactive allotments or resources where user run their applications manually
- Useful for:
 - Debugging Applications
 - Running GUI Applications

- **Batch Jobs**

- Non-interactive allotment of resources that run applications in the background
- Think “baking a batch of cookies”
- Useful for:
 - Applications that take a substantial amount of time
 - Non-interactive Applications

Scheduling an interactive job

- To work with R interactively, we request time from Summit
- When the resources become available the job starts
- Commands to run:

```
sinteractive --reservation=appm --time=00:10:00
```

*note “--reservation=appm” is only for this workshop

- Once we receive a prompt, then:

```
module load R  
R
```

- Once we finish we must exit! (job will time out eventually)

Scheduling Batch Jobs

- `sbatch` command: schedule a batch job with Slurm
- The `sbatch` command usually takes in 1 parameter: A job script.
 - Job scripts provide all information on what is needed for the job.
 - Parameters can be overwritten or added externally by specifying the parameter as a flag.
- Example:

```
sbatch test.sh
```

or

```
sbatch test.sh --time=02:00:00
```

Anatomy of a job script

```
#!/bin/bash
```

Job parameters

```
#SBATCH --ntasks=1           # Number of requested tasks
#SBATCH --time=0:01:00       # Max wall time
#SBATCH --partition=blanca-appm # Specify APPM nodes
#SBATCH --qos=blanca-appm-student # Specify you are student
#SBATCH --output=test_%j.out  # Rename standard output file
```

Your job commands

```
# Now run job (commands below here)

# Purge and load needed modules
module purge
module load R

# Run commands
Rscript myscript.R
```

Job Parameters

- Allocation: `--account=<account-name>`
- Partition: `--partition=<partition-name>`
- Number of nodes: `--nodes=<nodes>`
- Number of Tasks: `--ntasks=<number-of-tasks>`
- Quality of service: `--qos=<qos>`
- Reservation: `--reservation=<name>`
- Wall time: `--time=<wall-time>`
- Job Name: `--job-name=<jobname>`
- Output File: `--output=<outputname>`

More on slurm commands:
<https://slurm.schedmd.com/quickstart.html>

FYI: You do NOT actually type <> above – this designates something specific you as a user must enter about your job

APPM QoS, Account, and Partition

- When using the APPM Blanca nodes, job scripts should specify the following information:

```
--account=blanca-appm  
--partition=blanca-appm  
--qos=blanca-appm-student #for some in APPM qos may just be blanca-appm
```

- Users of Blanca also have access to a low-priority preemptible QoS that will attempt to run your job on any available node on Blanca
 - Preemptable jobs will be booted off the node and restarted at a later time if a high priority user runs a job on the node.

```
--qos=preemptable
```

schedule_hostname.sh

```
#!/bin/bash
#SBATCH --nodes=1                # Number of requested nodes
#SBATCH --ntasks=1              # Number of requested cores
#SBATCH --time=00:01:00         # Max wall time
#SBATCH --qos=blanca-appm-student # Specify QOS; may be just "blanca-appm" for some
#SBATCH --partition=blanca-appm  # Specify APPM nodes
#SBATCH --account=blanca-appm    # Specify account
#SBATCH --output=hostname_%j.out # Rename standard output file
#SBATCH --job-name=hostname      # Job name

# purge all existing modules
module purge

hostname
```

Running the job script

Schedule the job:

```
$ sbatch --reservation=appm schedule_hostname.sh  
(note “--reservation=appm” for this class only)
```

Check the status of the job:

```
$ squeue / $ squeue -u <user> /  
$ squeue -q <qos>  
...or  
$ sacct / $ sacct --format=<options>  
...or  
$ scontrol show job <job number>
```

Look at the job output:

```
$ cat hostname_<job-id>.out
```

*(*note that **<job-id>** is your job number)*

More on slurm commands: <https://slurm.schedmd.com/quickstart.html>

Topics we didn't cover today

- Job arrays (when you need to run lots of similar tasks)
- Running preemptable jobs
- CURC OnDemand (Interactive interface <https://ondemand.rc.colorado.edu>)
- Use of conda on CURC
- ..and lots of other cool stuff
- See: <https://curc.readthedocs.io> ; or email rc-help@colorado.edu to schedule a consultation

Thanks!

- Please fill out the survey: <http://tinyurl.com/curc-survey18>
- Contact: rc-help@Colorado.edu, Andrew.Monaghan@colorado.edu
- Course materials for today:
 - https://github.com/ResearchComputing/APPM_HPC
- Blanca (and other) documentation:
<https://curc.readthedocs.io/en/latest/access/blanca.html>
- Slurm Commands: <https://slurm.schedmd.com/quickstart.html>

Supplemental Slides

Environment variables

- Environment variables store important information needed by Linux users, programs, etc.
- Type '`env`' to see your currently set environment variables
- Useful Environment variables:
 - `PATH`: directories to search for commands
 - `HOME`: home directory
 - `PWD`: current working directory
 - `USER`: username
 - `LD_LIBRARY_PATH`: directories to search for shared objects (dynamically-loaded libs)

Shell Wildcards and Special Characters

- * - matches zero or more characters
- ? - matches a single character
- # - comment; rest of the line is ignored
- \ - escape; don't interpret the next character

Modes (aka permissions)

- Three classes of users:
 - User (u) aka “owner”
 - Group (g)
 - Other (o)
- Three types of permissions
 - Read (r)
 - Write (w)
 - Execute (x)

.. own grp oth
-|---|---|---

drwxr-xr--

Modes (continued)

- `chmod` changes modes:
- To add write and execute permission for your group:
`chmod g+wx filename`
- To remove execute permission for others:
`chmod o-x filename`
- To set only read and execute for your group and others:
`chmod go=rx filename`